

Machine Learning

Live Session #4

Review of Training Neural Networks

Neural Networks and Deep Learning

- Weights W are initialized randomly
 - Unless another initialization is specified (e.g., from a previous model or experience)
- Then, compare the predictions of the model with the actual labels in the training data
- Based on the comparison, update the weights to get the predictions closer to the actual labels
- Repeat the previous two steps until stopping criterion is met
- How do we make the comparison and update the weights to find the optimal values W^* ?

Neural Networks and Deep Learning

- How do we make the comparison?
 - This is the role of the loss function, which is *only* a function of the weights W , given the training data and architecture of the NN

$$Loss(W; \text{training set}, \text{architecture of NN})$$

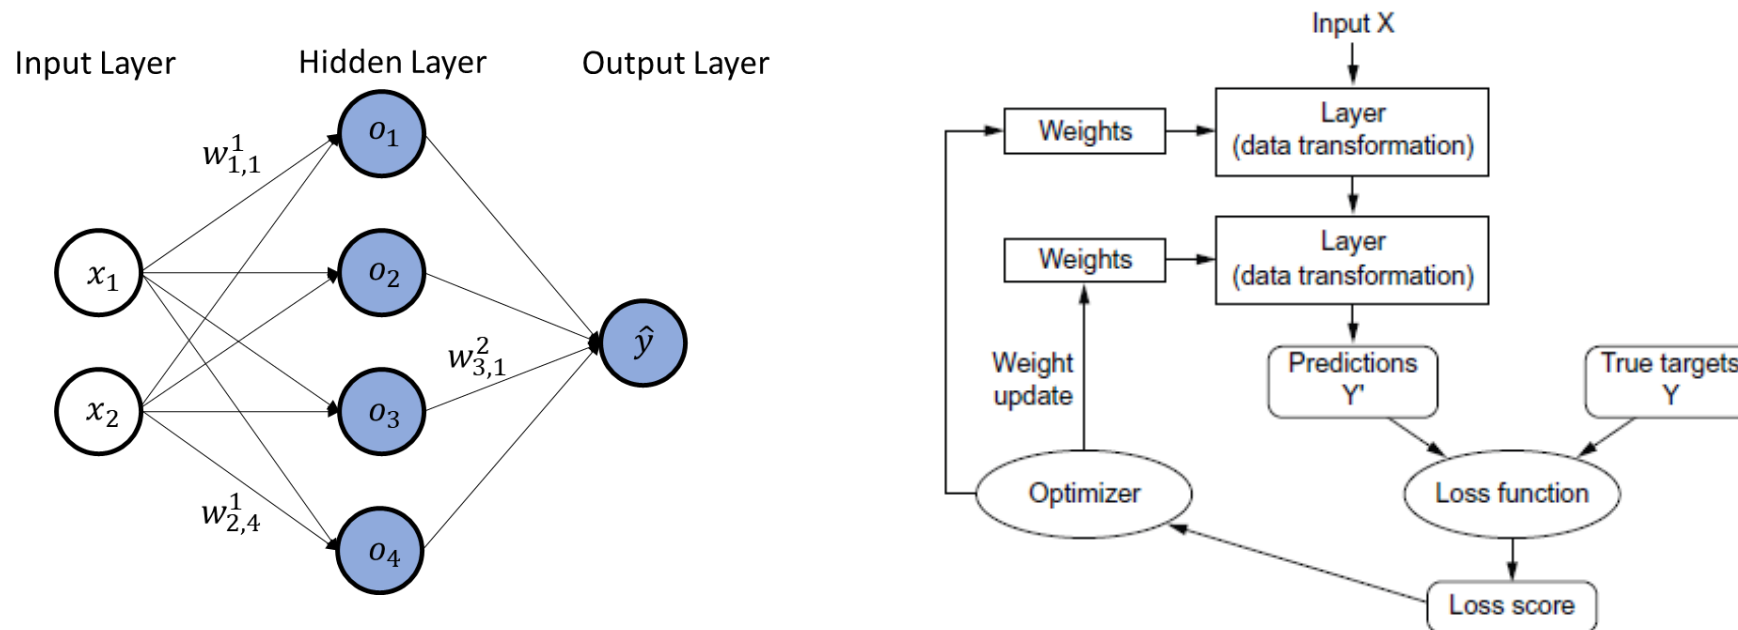
- Larger value of the loss means the compared values were further away
- Smaller value of the loss means the compared values were closer

Details of Training Process

- Important: for a particular training set and architecture, the loss is only a function of the weights W

$$\text{Loss}(W; \text{training set, architecture of NN})$$

- The loss function is used along with an optimization algorithm to determine how to change the weights to improve the loss
- The goal of the optimization is to find the optimal weight values W^* that minimize the loss



Neural Networks and Deep Learning

- To calculate the loss for a given set of weight values W
 1. Each x in the training set is put through the network to obtain a prediction $\hat{y}(x ; W)$
 2. Calculate the loss function (example below using mean squared loss function for regression)


$$Loss(W; \text{training set, architecture of NN}) = \frac{1}{n} \sum_{i=1}^n (\hat{y}(x_i ; W) - y_i)^2$$

x_1, x_2, \dots, x_n are training observations and y_1, y_2, \dots, y_n are actual labels

Training Dense Feed-Forward Neural Networks

Income (x_1)	Age (x_2)
22003	45
57230	54
75137	28
31208	54
54078	23
44413	44
55237	46

Input Layer Hidden Layer Output Layer



Computation Node

$$o_3 = f_1(\text{Weights} \times \text{Inputs} + \text{Bias})$$

$$o_3 = f_1(w_{1,3}^1 \times x_1 + w_{2,3}^1 \times x_2 + w_{0,3}^1)$$


Training Dense Feed-Forward Neural Networks

Income (x_1)	Age (x_2)
22003	45
57230	54
75137	28
31208	54
54078	23
44413	44
55237	46

Input Layer

Hidden Layer

Output Layer




Computation Node

$$\hat{y} = f_2(\text{Weights} * \text{Inputs} + \text{Bias})$$

$$\hat{y} = f_2(w_{1,1}^2 o_1 + w_{2,1}^2 o_2 + w_{3,1}^2 o_3 + w_{4,1}^2 o_4 + w_0^2)$$


Training Dense Feed-Forward Neural Networks



Training Dense Feed-Forward Neural Networks

**Update weights based
on loss function and
optimizer**

Income (x_1)	Age (x_2)
22003	45
57230	54
75137	28
31208	54
54078	23
44413	44
55237	46



Predicted Revenue Given Set of Weights W		Revenue from Customer
$\hat{y}(22003,45; W)$	The loss function compares these	14.03875
$\hat{y}(57230,54; W)$		23.31168
$\hat{y}(75137,28; W)$		24.05046
$\hat{y}(31208,54; W)$		18.5386
$\hat{y}(54078,23; W)$		18.50195
$\hat{y}(44413,44; W)$		20.63106
$\hat{y}(55237,46; W)$		22.32953

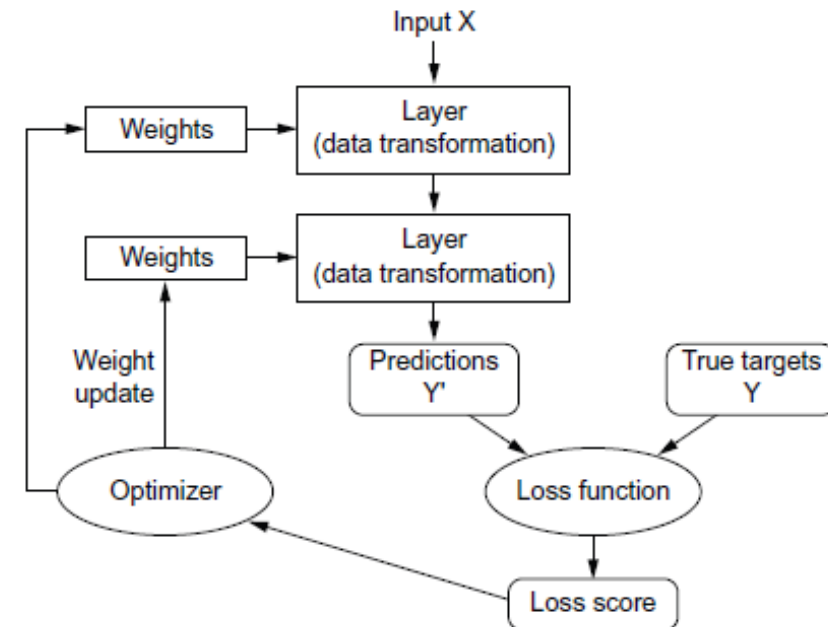
$Loss(W; \text{training set, architecture of NN})$

$$= \frac{1}{n} \sum_{i=1}^n (\hat{y}(x_i; W) - y_i)^2$$

Neural Networks and Deep Learning

- When the training set is too large, it's not feasible to calculate the loss based on the entire training set every iteration
 - Solution: use batches consisting of subsets of rows of training data
 - Take for example a batch size of 2

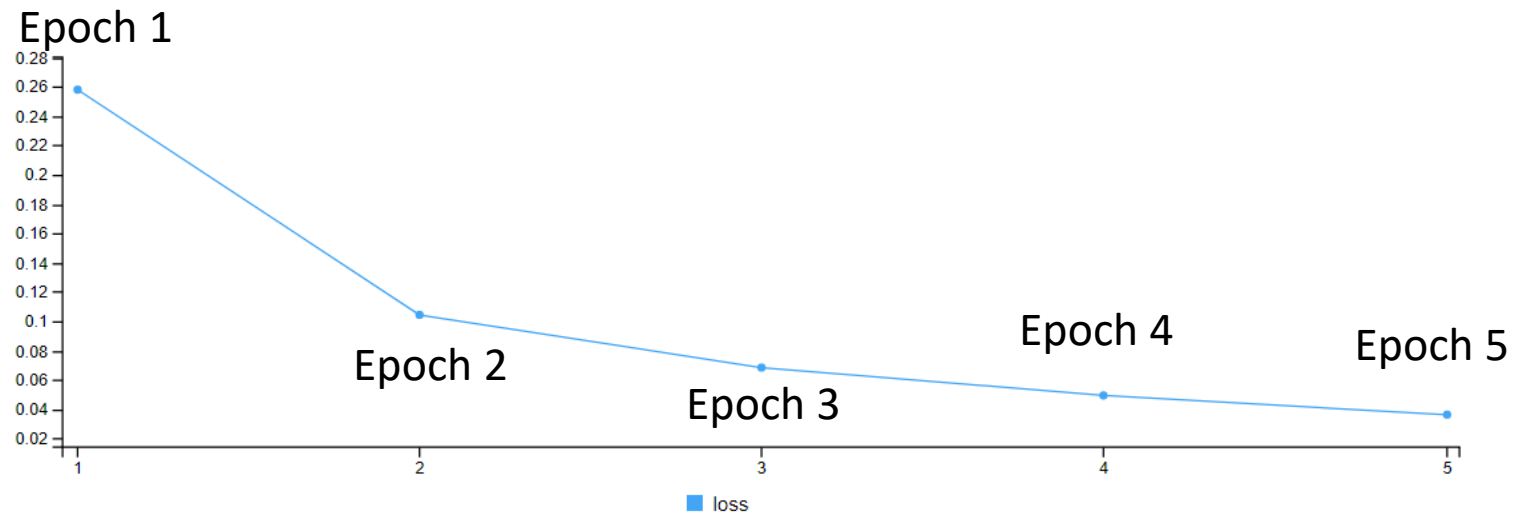
Income (Customer Feature 1)	Age (Customer Feature 2)	Revenue from Customer	Predicted Revenue Given Set of Weights W
22003	45	14.03875	$\hat{y}(22003,45; W)$
57230	54	23.31168	$\hat{y}(57230,54; W)$
75137	28	24.05046	$\hat{y}(75137,28; W)$
31208	54	18.5386	$\hat{y}(31208,54; W)$
54078	23	18.50195	$\hat{y}(54078,23; W)$
44413	44	20.63106	$\hat{y}(44413,44; W)$
55237	46	22.32953	$\hat{y}(55237,46; W)$



Neural Networks and Deep Learning

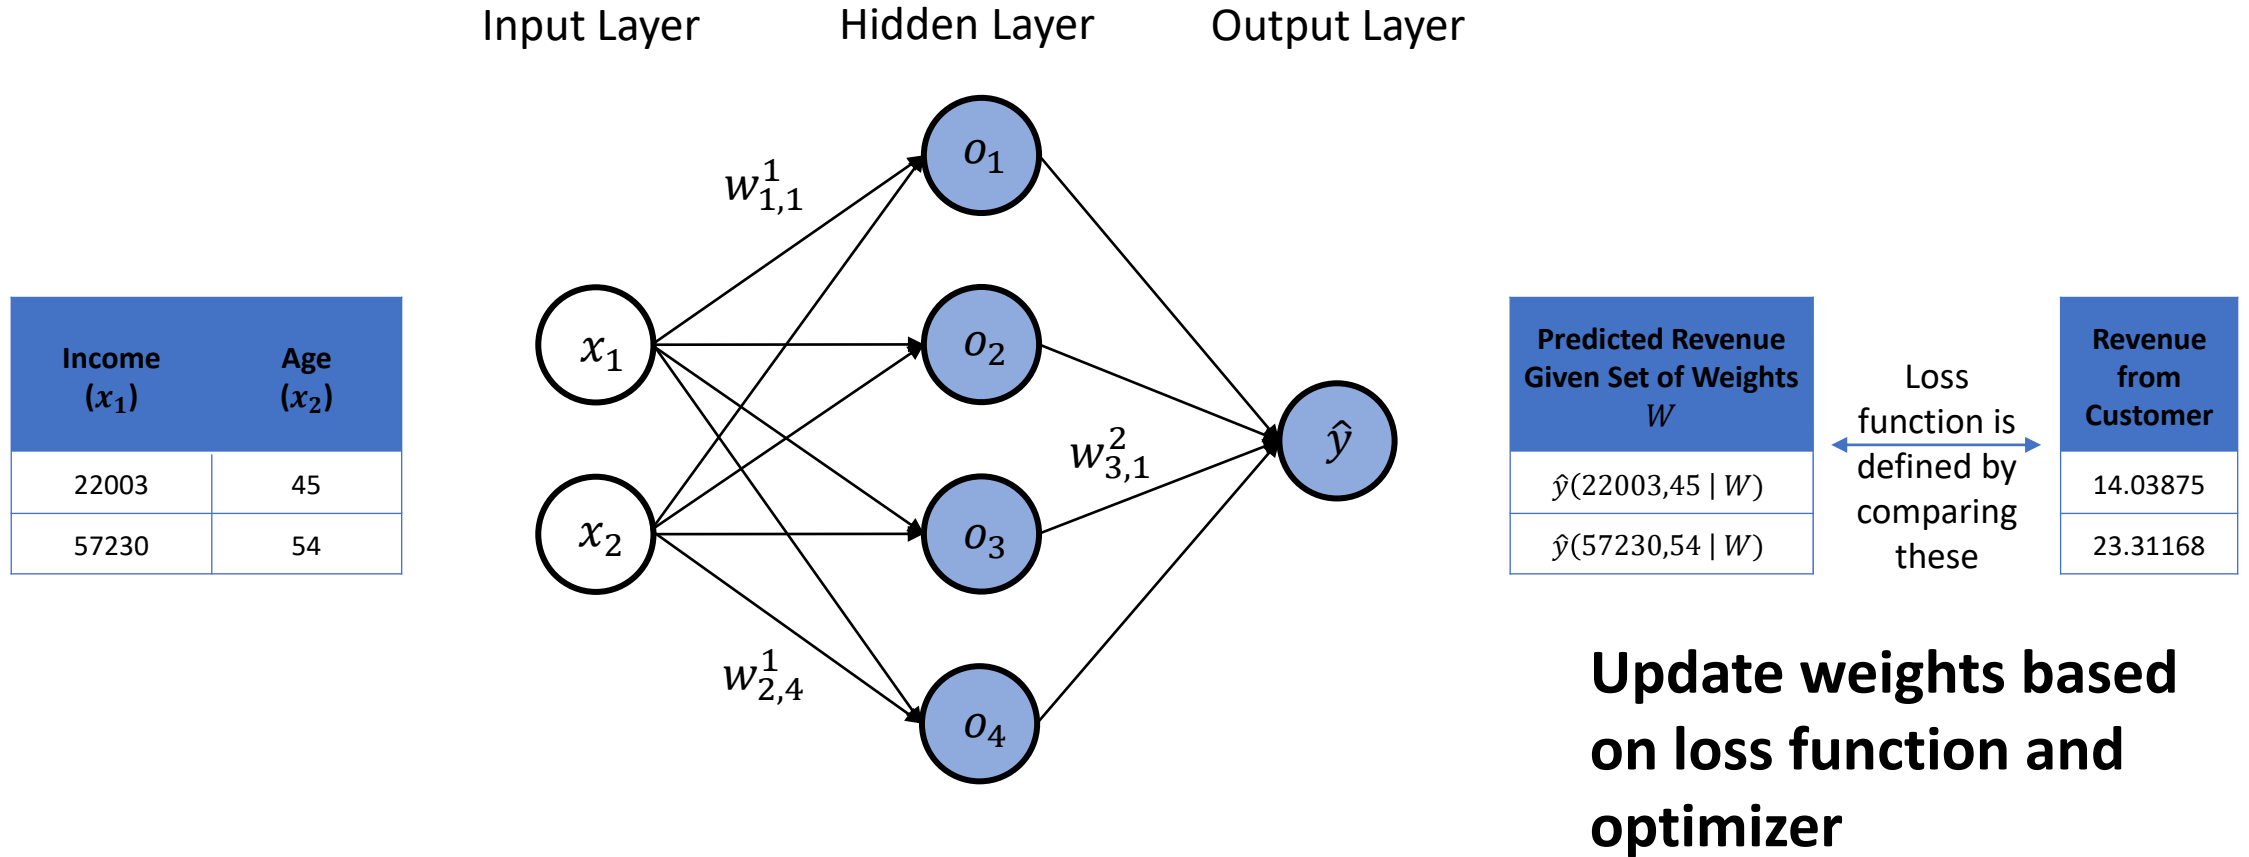
- Take for example a batch size of 2
 - Iterate through the training data taking two rows at a time
 - On each iteration, define the loss function
 $Loss(W|selected\ two\ rows\ of\ training\ data, architecture\ of\ NN)$
 - Use this loss function in conjunction with optimization algorithm to determine how to update weights
 - After iterating through the entire training set, we have completed an “epoch”
 - We specify the number of epochs in advance

Income (Customer Feature 1)	Age (Customer Feature 2)	Revenue from Customer	Predicted Revenue Given Set of Weights W
22003	45	14.03875	$\hat{y}(22003,45; W)$
57230	54	23.31168	$\hat{y}(57230,54; W)$
75137	28	24.05046	$\hat{y}(75137,28; W)$
31208	54	18.5386	$\hat{y}(31208,54; W)$
54078	23	18.50195	$\hat{y}(54078,23; W)$
44413	44	20.63106	$\hat{y}(44413,44; W)$
55237	46	22.32953	$\hat{y}(55237,46; W)$



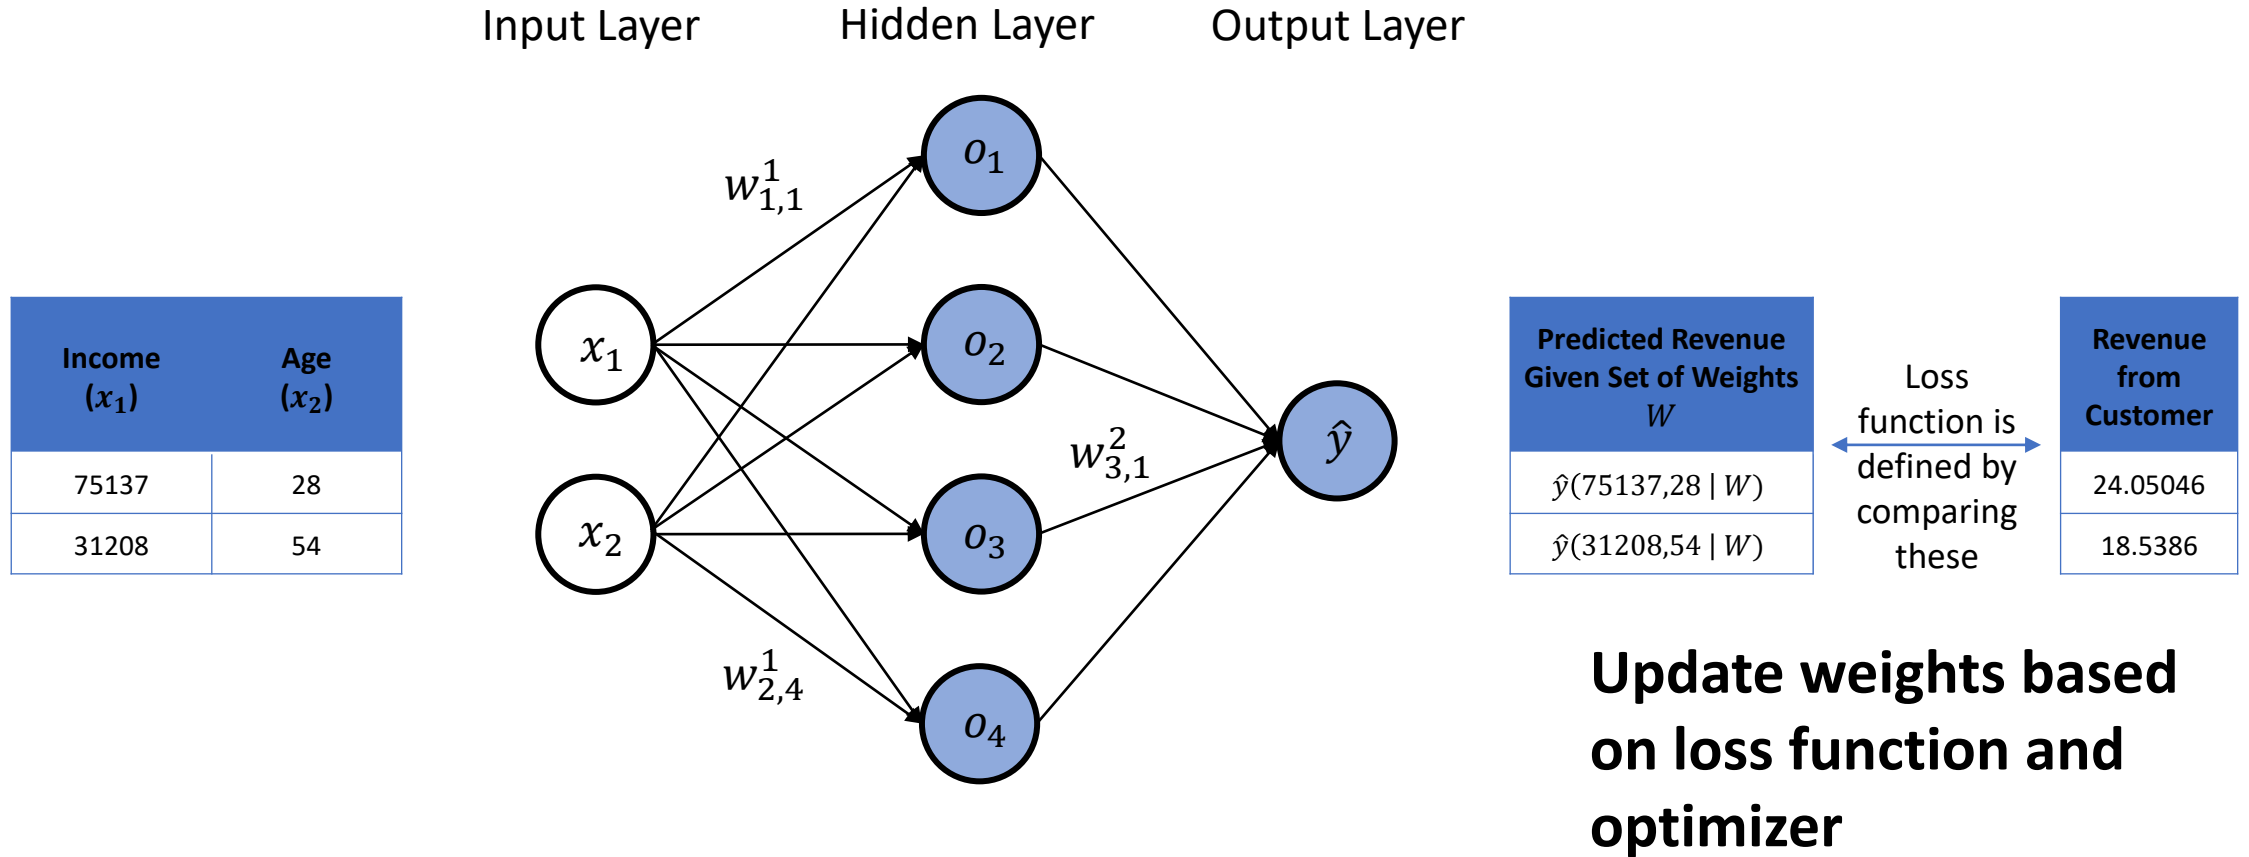
Training Dense Feed-Forward Neural Networks

Epoch #1, Step #1



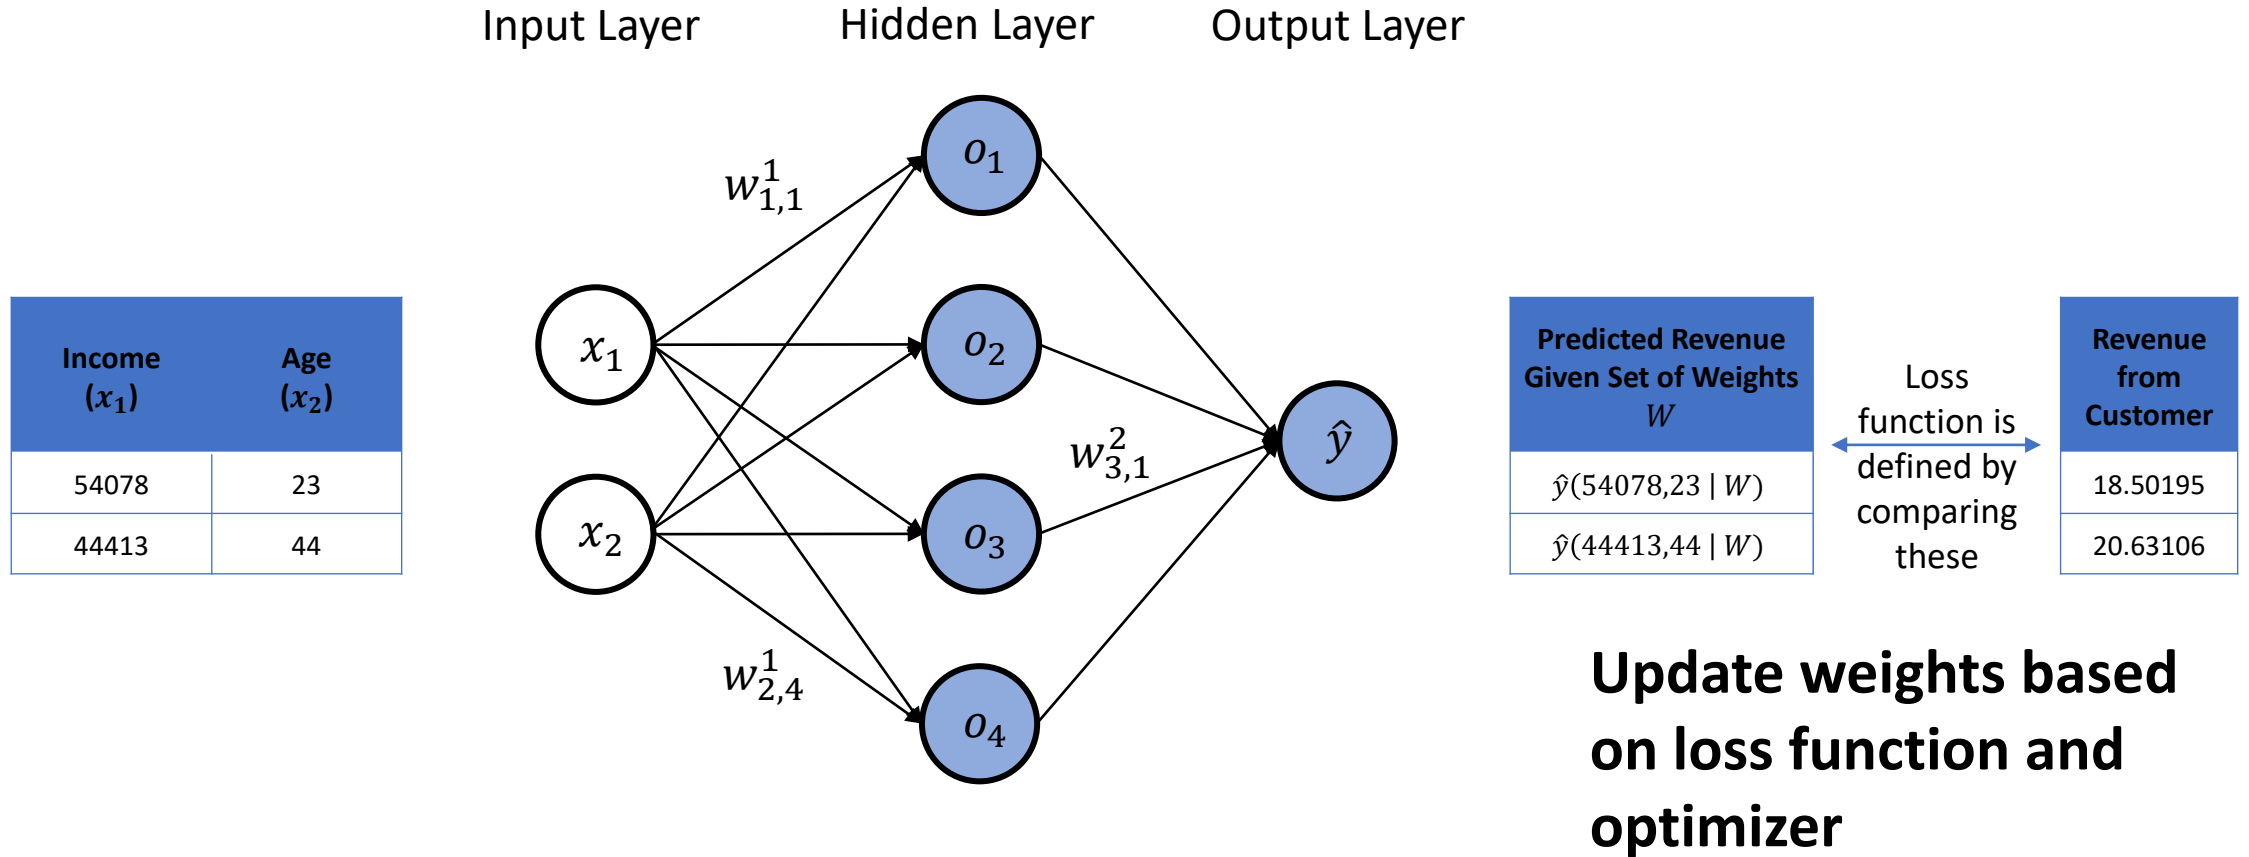
Training Dense Feed-Forward Neural Networks

Epoch #1, Step #2



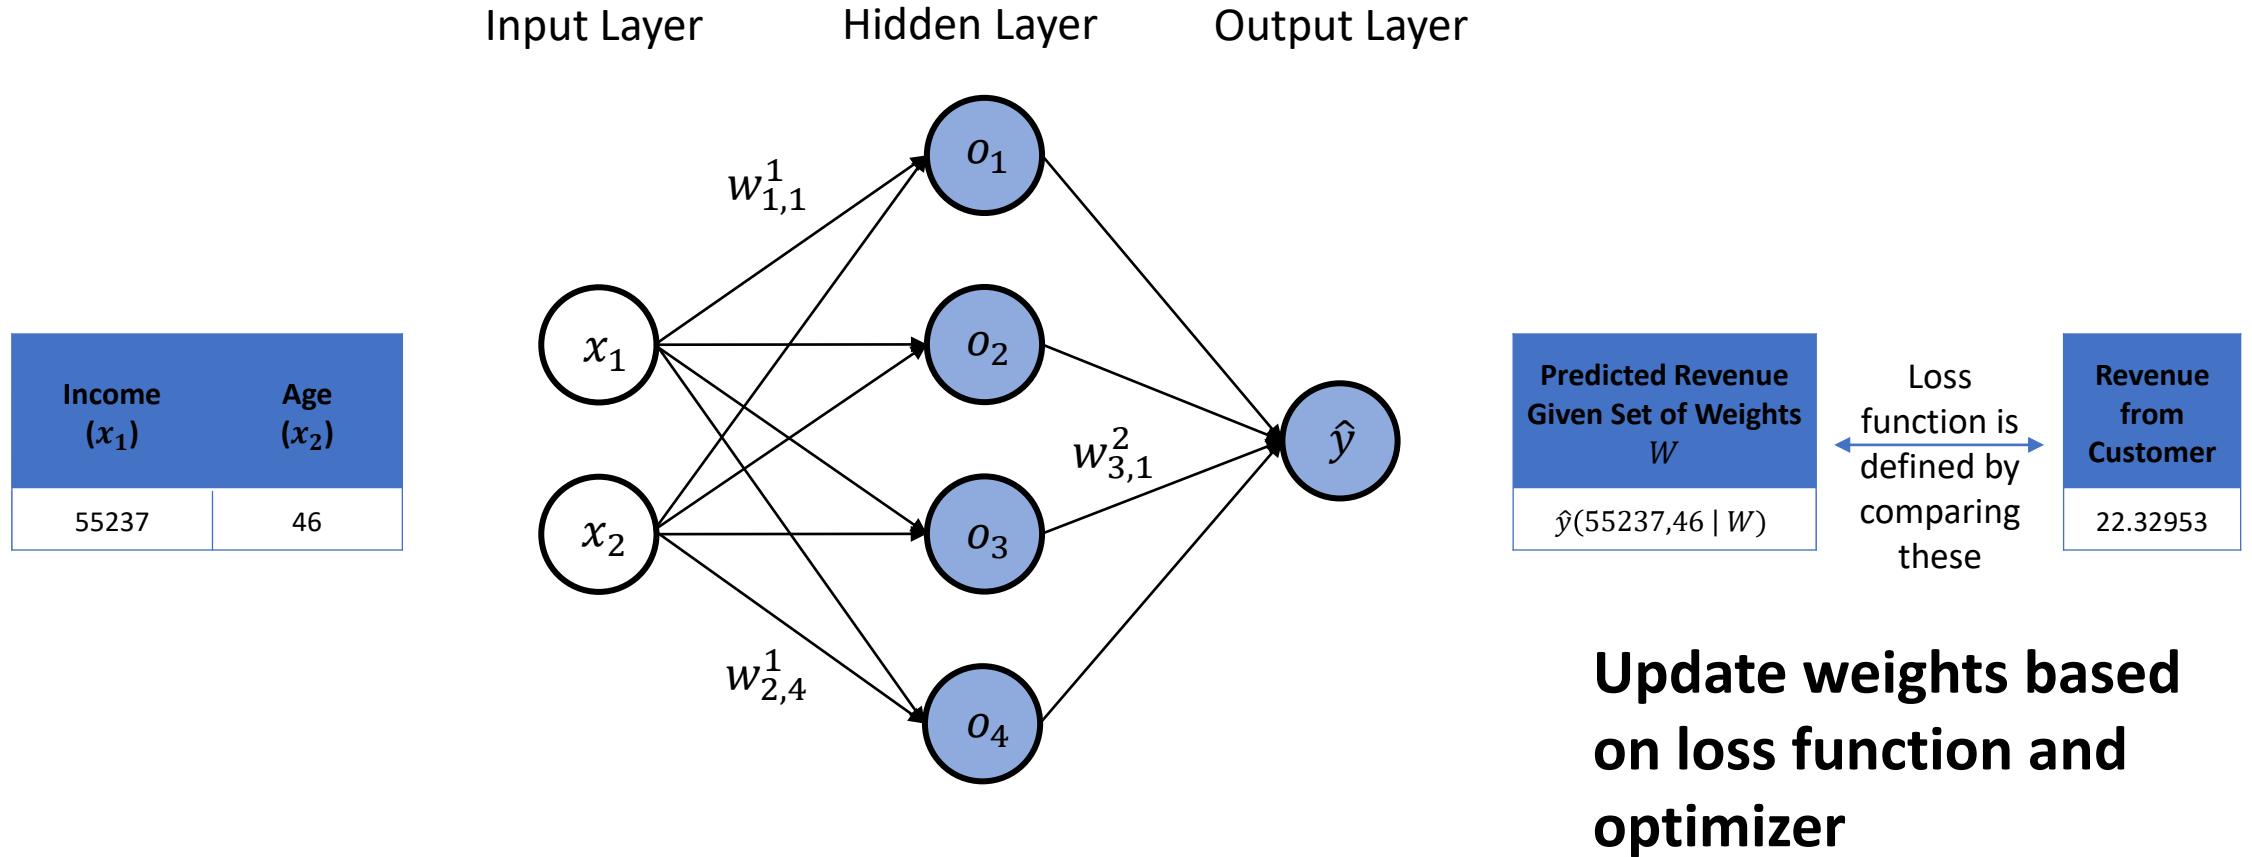
Training Dense Feed-Forward Neural Networks

Epoch #1, Step #3



Training Dense Feed-Forward Neural Networks

Epoch #1, Step #4



Training Dense Feed-Forward Neural Networks

- Finished iterating through the training set
 - Epoch #1 is now complete
- Shuffle the training set and start Epoch #2
- Continue in same manner until all epochs are completed

Income (Customer Feature 1)	Age (Customer Feature 2)	Revenue from Customer
22003	45	14.03875
57230	54	23.31168
75137	28	24.05046
31208	54	18.5386
54078	23	18.50195
44413	44	20.63106
55237	46	22.32953

